

AI in sociological research

Some ways I think we can use these tools
in our own research

Last year's discussion focused on student responses and essays.
Today I want to look at what happens inside a research project.

[Open the demonstration files](#)



Jacob Coxon 
@hilbertspaess

I resigned from Anthropic today. I spent the last three years doing pretraining research at both OpenAI and Anthropic. Neither company is acting responsibly. They are racing straight to self-improving superintelligence and gambling with our lives. More thoughts below.

8:04 PM · 9/8/26 · **812K** Views




590



2.8K



19K



6K





Home News Sport Business Technology Health Culture Arts Travel Earth Audio Video Live Documentaries

Anthropic boss Dario Amodei calls for AI development to slow down

 Summarize this page using Acrobat AI Assistant. 

3 days ago

Share 

Save 

 Add as preferred on Google

Paulin Kola and **Lily Jamali**, North America technology correspondent, in San Francisco



Trump rejected the idea of placing guardrails on AI, saying on Truth Social that the only control needed is “a STRONG AND SMART (High IQ!) PRESIDENT, and the U.S.A. has that, in spades!”

Cognitive offloading and learning

External resources change what we need to remember and do ourselves.

Trends in Cognitive Sciences

CellPress

Review

Cognitive Offloading

Evan F. Risko^{1,*} and Sam J. Gilbert^{2,*}

2016 · Review

Risko & Gilbert

Offloading reduces cognitive demands. Consequences depend on the task and use of the aid.

The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects From a Survey of Knowledge Workers

Hao-Ping (Hank) Lee

Carnegie Mellon University

Pittsburgh, Pennsylvania, USA

haopingl@cs.cmu.edu

Advait Sarkar

Microsoft Research

Cambridge, United Kingdom

advait@microsoft.com

Lev Tankelevitch

Microsoft Research

Cambridge, United Kingdom

levt@microsoft.com

Ian Drosos

Microsoft Research

Cambridge, United Kingdom

t-iandrosos@microsoft.com

Sean Rintel

Microsoft Research

Cambridge, United Kingdom

serintel@microsoft.com

Richard Banks

Microsoft Research Cambridge

Cambridge, United Kingdom

rbanks@microsoft.com

Nicholas Wilson

Microsoft Research

Cambridge, United Kingdom

niwilson@microsoft.com

2025 · Workplace survey

Lee et al.

319 workers: confidence in AI was associated with less reported critical thinking.

PNAS

RESEARCH ARTICLE

ECONOMIC SCIENCES

Generative AI without guardrails can harm learning: Evidence from high school mathematics

Hamsa Bastani^{a,b,1}, Osbert Bastani^{c,1}, Alp Sungu^{a,1,2}, Haosen Ge^b, Özge Kabakcı^d, and Rei Mariman^e

Affiliations are included on p. 7.

Edited by Emma Brunskill, Stanford University, Stanford, CA; received November 3, 2024; accepted May 5, 2025 by Editorial Board Member Mark Granovetter

2025 · Randomized school trial

Bastani et al.

Unaided exam grades fell after unrestricted AI practice help. A guided tutor largely removed the penalty.

Reduced effort, reduced learning and lasting cognitive loss are different claims.

[Actual paper excerpts · click a screenshot to open its source](#)

What research agents can accomplish

Article

Accelerating scientific discovery with Co-Scientist

<https://doi.org/10.1038/s41586-026-10644-y>

Received: 20 March 2025

Accepted: 11 May 2026

Published online: 19 May 2026

Open access

Check for updates

Juraj Gottweis^{1,8}, Wei-Hung Weng^{2,8}, Alexander Daryin^{1,8}, Tao Tu^{2,8}, Petar Sirkovic^{1,8}, Artiom Myaskovsky^{1,8}, Grzegorz Glowaty^{1,8}, Felix Weissenberger^{1,8}, Alessio Orlandi^{1,8}, Dan Popovici³, Anil Palepu³, Keran Rong², Ryutaro Tanno², Khaled Saab², Fan Zhang³, Jacob Blum⁴, Andrew Carroll⁵, Kavita Kulkarni⁵, Nenad Tomašev², Dina Zverinski¹, Ivor Rendulic¹, Elahe Vedadi², Florian Hasler¹, Luka Rimanic¹, Marina Boia¹, Ivan Budiselic¹, Ben Feinstein³, Mathias Bellaïche³, Tom Sheffer³, Jan Freyberg², Jeremy Ratcliff², Ottavia Bertolli², Katherine Chou³, Avinatan Hassidim³, Burak Gokturk¹, Amin Vahdat¹, Yuan Guan⁴, Vikram Dhillon⁵, Eeshit Dhaval Vaishnav⁶, Byron Lee⁶, Tiago R. D. Costa⁷, José R. Penadés⁷, Gary Peltz⁴, Yossi Matias³, James Manyika³, Demis Hassabis², Yunhan Xu², Pushmeet Kohli², Annalisa Pawlosky³, Alan Karthikesalingam² & Vivek Natarajan²

Google · Nature, 2026

Co-Scientist

Agents proposed hypotheses; scientists selected and tested them in the lab.

A system can extend what a researcher can execute while still needing their judgment.

[METR: what a time horizon measures](#)

Actual paper excerpts · selected studies, not a head-to-head comparison.

Can AI agents conduct open-ended AI research?

Early evidence from two case studies

Peter Kirgis^{1,*}, Sayash Kapoor^{1,*}, Andrew Schwartz^{1,2}, Stephan Rabanser¹, David Africa³, Konstantinos Voudouris³, Viet Nguyen⁴, Toby Pilditch³, Magda Dubois³, Harry Coppock³, Cozmin Ududec³, Nitya Nadgir⁵, Matilda Orona⁶, Tilman Bayer⁵, Derrick Chan-Sew⁵, Yue Ling⁵, Abhishek Shetty⁵, Helen Toner⁷, Gillian Hadfield⁸, Seth Lazar⁸, Steve Newman⁹, Shoshannah Tekofsky¹⁰, Rishi Bommasani¹¹, Arvind Narayanan¹

¹Princeton University ²Cornflower Labs ³UK AI Security Institute ⁴University of Toronto ⁵Independent ⁶UC Berkeley ⁷Georgetown University (CSET) ⁸Johns Hopkins University ⁹Golden Gate Institute for AI ¹⁰AI Digest ¹¹Stanford University

*Equal contribution †CRUX Core Team

Kirgis et al., including Narayanan · 2026 preprint

Two open-ended research tasks

Six days of research: engineering completed, but both papers judged below the publication bar.

2025: the chat window

Illustrated exchange

Summarize research on chatbots and support-seeking from friends and family.

A prose response appears here, with suggested themes and references to inspect.

Ask a follow-up...

The researcher carries the workflow

Read the answer.


Find the papers.

Move material into files.

Run analysis elsewhere.

The use case discussed last year. Tool-using and multi-agent systems already existed in 2025.

2026: working with a team of agents



The harness manages execution around the model. The researcher remains part of the revision loop.

[Open the four-stage interactive diagram](#)

Teaching schematic; no agents are running on the slide.

1. Complete cognitive offloading

- Give it a bounded task and inspect what comes back.
- “Reproduce this figure from the supplied data and document any discrepancies.”
- Other uses: bibliographic records, file conversion, replication checks.

The failure is accepting an output because it looks finished.

2. Guided collaboration

- Let it do much of the work; intervene where a decision matters.
- “Show me the proposed measures before running the analysis.”
- Other uses: developing an analysis, revising a coding scheme, comparing operationalizations.

The failure is approving each step without really examining it.

3. The pain-in-the-ass coauthor

You bring an argument. It questions it. You respond, and work through the implications together.

“I think these findings support displacement. Here is why.
What am I missing?”

Useful for theory, measurement, identification and interpretation.
Its objections need scrutiny too; it can misunderstand the evidence.

Do chatbots change how people seek support?

Heavy ChatGPT users tend to be more lonely, suggests research

Rachel Hall, The Guardian, 25 March 2025

Our question: do AI chatbots reduce requests for support from friends and family?

[Open the article](#)

What would you actually type?

Open a project folder, add the skill, then give it the question.

Codex CLI / IDE: `$literature-audit`

Claude Code: `/literature-audit`

“Find primary studies. Check the outcome instruments and source versions. Write an evidence table and bibliography. Ask me about consequential ambiguities.”

[Open the skill and setup instructions](#)

Workshop-authored skill; search and file tools depend on the session.

Inspect the evidence table and retrieval record

The agent assembled nine records. The useful part is seeing what each study actually measured.

Study design	What was assigned, and what was self-selected?
Outcome	Requests, socialization, loneliness, or response ratings?
Source check	Which version? Full text or abstract only?

[Open evidence table](#)

[Open retrieval record](#)

[Open Fang et al., v2](#)

Working through a measurement decision

RESEARCHER

“Could we compare the share of support needs that lead to a friend?”

AI COAUTHOR

“What if chatbot use changes what people report as a need?
Then the denominator changes too. Could total requests over
a fixed follow-up be the primary outcome?”

[Open the full exchange](#)

Scripted teaching example; the actual reviewer reports are linked separately.

Two agents checked the proposal

- The source reviewer checked the claim against the paper.
- The methods reviewer examined the outcome, comparison and analysis.
- They then read each other's objections. One unresolved issue: a contact-prompt effect does not establish a bridging mechanism.

[Open the actual reports and cross-checks](#)

The resulting research proposal

A short protocol with an evidence assessment, defined outcome, assignment scheme, analysis and draft diary questions.

- Compare a standard chatbot offer with usual access.
- Separately compare an added human-contact feature with the standard offer.
- Retain non-users in assigned groups; missing diaries are not zero requests.

[Open the compiled PDF](#)

[Open LaTeX](#)

[Open bibliography](#)